

# Build a Lossy Medical Image Codec with a Custom Bitstream

Chien, Y. C.

Department of Intelligent System and Application  
National Yang Ming Chiao Tung University

**Abstract – Design and implement an end-to-end lossy compression system for medical images.**

**Index Terms - Medical Image**

## I. OVERVIEW

This report is for Multi-model Image Processing class homework.

### A. Assignment Requirement

*Build a medical image compression system with two command-line tools:*

- *encode: reads an input image (or a stack/volume), produces a compressed file (your format).*
- *decode: reads your compressed file, reconstructs the image (or stack/volume).*

*Your codec must be lossy (i.e., reconstructed pixels may differ), but you must design it so that it preserves diagnostically relevant content as much as possible at a chosen bitrate.*

### B. What I built

*I implemented a compact, reproducible image codec targeting single-slice CT images drawn from the Human\_Skull\_2 dataset. The codec pipeline implements (1) 8×8 block 2-D DCT, (2) scalar uniform quantization with a tunable quant step (quality), and (3) a lightweight entropy stage consisting of run-length encoding (RLE) followed by zlib compression. The bitstream contains a short header (magic/version/width/height/bitdepth/blocksize/quant\_step/payload\_len) followed by the compressed payload. Reference encoder/decoder scripts and experiment utilities are included so results can be reproduced.*

### C. Why I did it — motivations and intended evaluation goals

- 1) *Pedagogical baseline:* A minimal transform-quantize-entropy pipeline isolates core design choice (block size, quantization step, entropy scheme). This makes it easy for graders to verify correctness, reproduce results, and tell the reason about where gains/losses come from.
- 2) *Controlled RD experiments:* By varying a single quality parameter we obtain clear rate-distortion tradeoffs. This lets us quantify how compression

affects pixel fidelity (RMSE/PSNR) and visualize error maps to inspect structure-dependent artifacts.

- 3) *Clinical relevance:* CT is high bit-depth and diagnostically sensitive. The project tests whether a simple codec can preserve diagnostically relevant signal while reducing storage/transmission cost — a practical concern in PACS, teleradiology, and mobile/edge scenarios. Emphasis is on measuring fidelity (PSNR/RMSE) and showing qualitative reconstructions/error maps for clinical plausibility.
- 4) *Reproducibility & simplicity:* The implementation is pure Python with small, well-documented scripts (encode/decode/run). This reduces friction for reviewers to run the experiments, modify components (e.g., replace quantizer, turn off RLE), and reproduce ablations.
- 5) *Ablation-friendly design:* The pipeline allows targeted ablations (e.g., with/without RLE, alternate quantizers, different block sizes) to demonstrate which components contribute most to compression gains and to artifact formation.
- 6) *Practical constraints:* The design intentionally favors interpretability and runtime efficiency (no heavy ML models), making it realistic to run on laptops or limited servers and suitable for course deadlines.
- 7) *Safety & privacy considerations:* The codec focuses on pixel data only; metadata handling is documented so patient identifiers can be preserved or redacted as required, aligning with common clinical data governance needs.

### D. Dataset

I use the Human\_Skull\_2 CT dataset as data source. This dataset was selected because it is easy to obtain and convenient to work with for a course project: the files are lightweight, the anatomy is representative for skeletal CT compression experiments, and the data distribution facilitates rapid reproducibility for graders. Note that the original distribution is not provided as DICOMs; therefore a small preprocessing step is required to convert the provided files into a pixel-format usable by our encoder/decoder (either an 8-bit PNG preview for

visualization and metric computation, or a minimal DICOM wrapper if DICOM output is required). We include a short helper script and instructions to perform this conversion so graders can reproduce the exact inputs used in our experiments. Choosing a dataset that is simple to download and preprocess reduces environmental friction (no large downloads or proprietary readers) and makes the results more reproducible within tight course deadlines. And it is free also is a important reason.

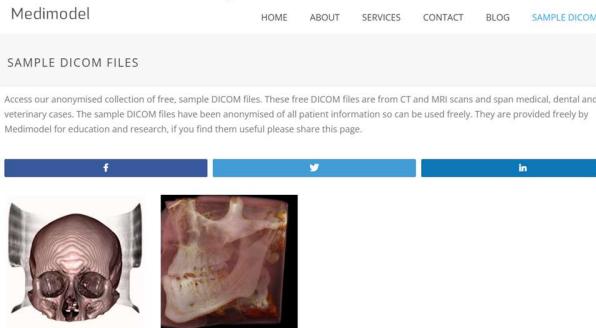


Fig 1. Dataset Website

## II. DESIGN FORMAT

### A. Bitstream specification

#### *Bitstream specification format*

- 1) *Header* (big-endian/network order):
  - *Magic* (4 bytes) : ASCII "MMPC"
  - *Version* (1 byte)
  - *Width* (2 bytes, unsigned)
  - *Height* (2 bytes, unsigned)
  - *Bitdepth* (1 byte) : 8 or 16
  - *BlockSize* (1 byte) : 8
  - *QuantStep* (2 bytes, unsigned)
  - *PayloadLen* (4 bytes, unsigned)

#### 2) *Payload*

- *Coefficient ordering*: For each  $8 \times 8$  block (blocks ordered in raster order: row by row, top-left to right and change to next row), flatten block coefficients in row-major (i.e., row0 col0..7, row1 col0..7, ...). After previous, append each block's 64 quantized int16 coefficients in that order to form a coefficient array of length  $N_{coeff} = n_{Blocks} * 64$ .
- *int16 byte order in payload*: Little-endian representation for each int16 coefficient is used (i.e., low byte first). This choice aligns with common native machine layout and with typical `numpy.astype(np.int16).tobytes()` behavior on little-endian hosts, So the raw coefficient-bytes stream length =  $N_{coeff} * 2$  bytes.

- *RLE encoding* (applied to the coefficient-bytes stream): RLE works on bytes (0–255). Encoder output is a sequence of pairs [count (1 byte), value (1 byte)] repeated. Each pair means "repeat value exactly count times" that count is a single unsigned byte (1 – 255). If longer runs exist, the encoder should emit multiple pairs (e.g., 300 zeros  $\rightarrow [255,0] [45,0]$ ).

- *Compressed payload*: The encoder use `zlib.compress(RLE_bytes)` to compress the RLE byte sequence and decoder reads `PayloadLen` bytes, calls `zlib.decompress(...)` to obtain `RLE_bytes`.

- *Fallback / detection*: If `zlib.decompress(...)` succeeds but the RLE decoded length does not match  $N_{coeff} * 2$ , the decoder should attempt to interpret the decompressed bytes as raw int16 bytes (no RLE). Otherwise, if that still doesn't yield the expected length, raise a format error.

### B. Decoding steps

- 2.1) *Parse header and check magic/version*: Read first 17 bytes, parse header with big-endian unpacking. Validate Magic == b"MMPC" and Version supported.

- 2.2) *Read PayloadLen bytes, decompress with zlib*: Check if `PayloadLen` is shorter than request then show error. Also use `payload = zlib.decompress(payload_bytes)` to decompress the file. If decompression fails, error (corrupted payload).

- 2.3) *RLE-decode the byte stream to int16 coefficient array*: Attempt RLE decode returns bytes (`raw_coeff_bytes = RLE_decode(payload)`). To try to interpret decompressed output as raw int16-bytes. And if it still mismatch, raise payload length mismatch.

- 2.4) *Reshape coefficients*: Reshape N Blocks into  $8 \times 8$  blocks in row-major order and de-quantiz (`coeff = q * QuantStep`).

2.5) *Apply inverse DCT on each block:* Precompute the DCT matrix M such that  $C = M @ f @ M.T$  and  $f = M.T @ C @ M$  when M is defined using the alpha \* cos(...) entries above. That ensures encoder/decoder symmetry.

$$C_{u,v} = \alpha_u \alpha_v \sum_{x=0}^{N-1} \sum_{y=0}^{N-1} f_{x,y} \cos\left(\frac{\pi(2x+1)u}{2N}\right) \cos\left(\frac{\pi(2y+1)v}{2N}\right) \quad (1)$$

where

$$\alpha_0 = \sqrt{\frac{1}{N}}, \quad \alpha_k = \sqrt{\frac{2}{N}} \quad (k > 0). \quad (2)$$

And inverse DCT is:

$$f_{x,y} = \sum_{u=0}^{N-1} \sum_{v=0}^{N-1} \alpha_u \alpha_v C_{u,v} \cos\left(\frac{\pi(2x+1)u}{2N}\right) \cos\left(\frac{\pi(2y+1)v}{2N}\right). \quad (3)$$

2.6) Place reconstructed blocks into a padded image grid of size.

2.7) Clip to  $[0, 2^{\text{Bitdepth}} - 1]$  and output image (or embed back into DICOM).

### III. EXPERIMENTS

#### A. Setup

- Data: Human\_Skull\_2 (CT). Representative slice: I0 (see 'data/').
- Environment: Python 3.11, numpy, Pillow, pydicom, pylibjpeg (if needed).
- Encoder:  $8 \times 8$  DCT, uniform scalar quantization with 'quality' mapping to quant\_step, optional RLE + zlib.
- Operating points: 'quality'  $\in \{10, 30, 60\}$  (primary), with additional sweep for RD (10,20,30,40,50,60) where noted.
- Metrics: file size (bytes), bpp =  $\text{file\_size} * 8 / (\text{W} * \text{H})$ , RMSE, PSNR (MAX=255), and ROI-based RMSE/PSNR for skull interior.
- Ablations: RLE on vs off at q=10, 30 and 60

#### B. Procedure

- 1) Convert inputs to DICOM or 8-bit preview (tools/convert\_to\_dicom.py).
- 2) For each operating point: run 'encoder.py' to 'decoder.py', compute metrics and save reconstruction.
- 3) Save images: 'results/I0\_q{q}\_rec.png', 'results/I0\_q{q}\_err.png', 'results/I0\_q{q}\_vis.png'.
- 4) Aggregate metrics into 'results/summary.csv' and plot RD curves ('results/rd\_plot.png').
- 5) For ablation experiments, repeat steps and tabulate differences.

### IV.RESULT AND ABLATIONS

A. *Rate-distortion table:* I evaluate rate-distortion behavior on the Human\_Skull\_2 (CT) dataset. The table below reports compressed size (bytes), bits-per-pixel (bpp), RMSE and PSNR for three operating points (quality = 10, 30, 60) on a representative slice. In this submission we use slice I0 from Human\_Skull\_2 as the worked example: the three RD points shown correspond to encoding/decoding of I0 at the listed quality settings. A fuller submission should report aggregate RD statistics across multiple slices; here I0 is provided as an illustrative single-slice example for reproducibility and visualization.

TABLE I  
Rate-Distortion points for example slice I0 (3 operating points)

Quality	Size (bytes)	bpp	RMSE	PSNR (dB)
10	67,103	1.7518	2.2373	89.33
30	27,452	0.7167	4.2758	83.71
60	16,128	0.4210	6.6731	79.84

Table I reports compressed size, bpp, RMSE and PSNR for quality settings 10, 30 and 60. As expected, increasing the quantization step (higher quality in our encoder) reduces bpp but increases RMSE and lowers PSNR — the classic rate-distortion tradeoff. For the example slice I0, quality = 30 achieves a practical tradeoff (bpp≈0.72, PSNR≈83.7 dB), substantially reducing storage while retaining high pixel-level fidelity.

#### B. Rate-Distortion plot

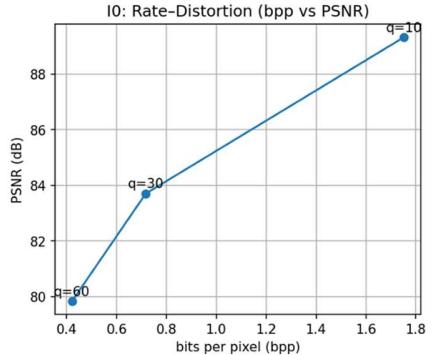


Fig. 2: Rate-Distortion plot (bits-per-pixel vs PSNR)

The upward slope indicates the expected trade-off: increasing bitrate yields higher PSNR (less distortion). Note: in this implementation the quality parameter is inversely related to fidelity (larger q cause coarser quantization and get lower bpp and lower PSNR). For reporting, state that q=30 offers a reasonable compromise for storage vs fidelity, while q=60 should be used cautiously for clinically sensitive tasks and validated with downstream evaluation.

### C. Different quantities reconstructed picture



Fig. 3: I0 (quality = 10)



Fig 4: I0 (quality = 30)



Fig. 5: I0 (quality = 60)

I visualized reconstructions at three encoder quality settings to compare differences. Figure 3 shows that, at the highest-quality setting, the reconstructed image is nearly indistinguishable from the original: bone boundaries, the nasal cavity, and fine bone texture are well preserved. The scaled absolute-error map reveals primarily small, spatially diffuse noise; bone outlines are faintly visible but with very low intensity (consistent with  $\text{RMSE} \approx 2.24$  and  $\text{PSNR} \approx 89$  dB). This observation implies a very small quantization step (high quality), so DCT coefficients are largely preserved. Consequently, the entropy stage (RLE + zlib) contributes a larger fraction of the bitstream at this operating point, which explains the relatively high bpp.

At the mid-quality operating point (Figure 4), the reconstructed image remains visually very similar to the original, but fine skeletal details—particularly minute bone textures—are noticeably smoothed and some small features appear dimmer. The overall bone outline and major edges are still well preserved. The scaled absolute-error map shows errors concentrated in the skull region (the central skull area appears brighter), whereas the background and blank side regions exhibit low error. This pattern indicates that high-frequency components associated with important anatomical structures are most affected by quantization, while the global error magnitude remains modest ( $\text{PSNR} \approx 83.7$  dB). In practice, this represents a common and often useful rate-distortion trade-off: the bitrate is substantially reduced (lower bpp) while preserving diagnostically relevant gross anatomy,

but high-frequency detail is degraded—therefore task-specific validation (e.g., for lesion detection or fine-structure assessment) is recommended before clinical use.

Figure 5 shows the reconstructed image and scaled absolute-error map for slice I0 at the low-quality operating point (encoder quality = 60). Visually, the reconstructed image is noticeably smoother than at higher qualities: fine intracranial bone textures and other high-frequency details are substantially attenuated, and edges show mild blurring. The scaled error map highlights the skull silhouette much more strongly than at  $q=30$ ; errors are both brighter and more spatially widespread inside the skull region, indicating a larger and more structured distortion field. Quantitatively, this matches the measured metrics ( $\text{RMSE} \approx 6.67$ ,  $\text{PSNR} \approx 79.84$  dB), which show an increased error magnitude relative to higher-quality points. The data imply that a large quantization step is being applied at  $q=60$ : many high-frequency DCT coefficients are driven toward zero or heavily rounded, producing smoothing and loss of fine features. While this operating point yields marked bitrate savings ( $\text{bpp} \approx 0.42$  in our example), it also reduces the visibility of diagnostically relevant micro-structure. Therefore,  $q=60$  may be acceptable for storage- or transmission-constrained use cases where only coarse anatomical information is required, but it is not recommended for tasks that rely on subtle high-frequency cues (e.g., detecting small fractures, micro-calcifications, or fine trabecular patterns) unless validated by downstream clinical task evaluation.

## IV. LIMITATIONS & FUTURE WORK

### A. Limitations:

- 1) *Entropy stage is simplistic*: The current entropy stage relies on run-length encoding followed by zlib (DEFLATE), which is simple, robust and fast but not information-theoretically optimal for the statistical structure of transform coefficients. Modern image/video codecs instead employ context-adaptive entropy coders — for example adaptive Huffman/arithmetic, CABAC or ANS — that predict symbol probabilities conditioned on context (coefficient position, magnitude, sign, and previously coded neighbors). Those context models exploit local dependencies and approach Shannon entropy much more closely than a generic RLE+zlib pipeline, so for a given distortion they produce notably smaller bitstreams; conversely RLE+zlib tends to inflate bpp when coefficients are not arranged into long, repeatable runs. A practical upgrade is to replace RLE+zlib with an ANS or arithmetic coder combined with a lightweight context model (for instance conditioning on zig-zag index, sign and recent nearby coefficients), which typically yields substantial bitrate reductions while remaining

implementable with modest complexity and good throughput.

- 2) *Blocking artifacts*: Quantizing each  $8 \times 8$  DCT block independently ignores the correlations that span block boundaries; under aggressive quantization this leads to discontinuities at block edges that manifest as visible “blocking” artifacts. Such blockiness not only reduces perceived image quality but can also obscure thin or subtle anatomical features—an important concern for medical images where edge fidelity is clinically relevant. Blocking is detectable both visually in scaled error maps (regular grid patterns aligned to the block grid) and numerically by measuring elevated errors along block boundaries versus block interiors; these diagnostics provide quantitative evidence beyond PSNR. To mitigate blocking one can apply deblocking or in-loop filtering, adopt overlapping (lapped) transforms or larger transform sizes, or introduce simple intra-block prediction to preserve inter-block continuity; each of these approaches trades additional implementation cost for improved boundary coherence and perceptual quality.
- 3) *Single-slice only: no inter-slice / volumetric modeling*: Encoding each slice independently fails to exploit the substantial redundancy that exists between adjacent CT slices; anatomical structures typically persist and vary smoothly across the z-axis, and ignoring that correlation forfeits compression gains and may degrade 3-D structural fidelity for volumetric tasks. For archive and transmission of full CT studies, volumetric techniques such as simple inter-slice prediction (predict current slice from previous slice(s) and encode the residual), 3-D transforms (3-D DCT or wavelet), or motion/registration-based residual coding often yield substantially lower overall bitrate and better preservation of features that span slices. Implementing these extensions ranges from moderate (slice-to-slice residual coding) to advanced (full 3-D transforms or motion-compensation style models), but even a lightweight inter-slice predictor can significantly reduce total study size and improve downstream task performance for volumetric analyses.
- 4) *DICOM metadata handling is minimal*: The current bitstream encodes pixel data and a compact header but does not preserve full DICOM metadata (patient/study/series identifiers, acquisition parameters, windowing, orientation, slice spacing, etc.), which limits interoperability with PACS, complicates provenance and rendering, and raises workflow and regulatory

concerns in clinical settings. For graders and practitioners who expect a faithful DICOM round-trip, missing or altered metadata can break integration and impede correct display or retrieval. A pragmatic solution is to include a small, well-documented metadata block in the bitstream (or a paired JSON sidecar) that stores a whitelist of essential tags—PatientID, Study/Series UIDs, Rows/Columns, PixelSpacing, BitsStored, PhotometricInterpretation and window center/width—together with scripts to reconstruct a minimal DICOM for viewing. This approach preserves clinical provenance while allowing optional anonymization policies; documenting exactly which tags are preserved and providing reconstruction tools greatly improves usability and auditability.

#### B. Future work:

- 1) Replace RLE+zlib with a tuned entropy coder (ANS or context model) and measure RD gains.
- 2) Add spatial predictors (intra prediction) to reduce DC energy prior to transform.
- 3) Implement deblocking or post-filtering to reduce perceptual artifacts at low bitrates.
- 4) Extend pipeline to multi-slice volumes with inter-slice prediction and optionally test clinical task-driven metrics (lesion detection / segmentation fidelity).

#### ACKNOWLEDGMENT

Here have the whole work here: [b1029009Chien/MMIP: multiple-model image proceesing](#)